

Crown-Jewel v0: visual (shift) and textual (paraphrase) invariance are BOTH saturated across zero-shot CLIP encoders, yet adversarial robustness is a separate axis that η/L_1 predicts and neither consistency does (ImageNet-100, per-tower points)

