


Crown-Jewel HARDENED: a 2x2 (modality x perturbation) view on zero-shot CLIP. Visual (shift) and textual (paraphrase) invariance are both saturated across 17 encoders. Only the image-adversarial cell (S) spreads across encoders and is predicted by η/L_1 ; the text worst-case S_{text} stays narrow, and η/L_1 is orthogonal to it once clean accuracy is controlled.

A. Both invariance cells saturate; the image-adversarial cell (S) spreads wide, the text cell (S_{text}) stays narrow



B. η/L_1 predicts the image cell (S), not the text cell (S_{text}); the consistencies predict neither

